

# Near-real time processing using Spark

## RDDs, DataFrame/DataSet, Spark Streaming, and Structured Streaming

Vladimir Sivčević

Senior Data Engineer

[www.vladsiv.com](http://www.vladsiv.com)

May 9, 2024

# Table of Contents

- **RDDs**

# Table of Contents

- RDDs
- Spark Streaming

# Table of Contents

- RDDs
- Spark Streaming
- DataFrame - DataSet

# Table of Contents

- RDDs
- Spark Streaming
- DataFrame - DataSet
- Structured Streaming

# RDD

## RDD - Resilient Distributed Dataset

# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

# RDD

RDD - Resilient Distributed Dataset

- **Resilient**
  - **Immutable** - Rebuild lost data



# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

- **Immutable** - Rebuild lost data
- **Fault Tolerant** - Recover data using Lineage

# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

- **Immutable** - Rebuild lost data
- **Fault Tolerant** - Recover data using Lineage

- **Distributed** - Across nodes, operated on in parallel

# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

- **Immutable** - Rebuild lost data
- **Fault Tolerant** - Recover data using Lineage

- **Distributed** - Across nodes, operated on in parallel

- Two ways to create them

# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

- **Immutable** - Rebuild lost data
- **Fault Tolerant** - Recover data using Lineage

- **Distributed** - Across nodes, operated on in parallel

- Two ways to create them

- *Parallelized Collections*

```
data = [1, 2, 3, 4, 5]  
distData = sc.parallelize(data)
```

# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

- **Immutable** - Rebuild lost data
- **Fault Tolerant** - Recover data using Lineage

- **Distributed** - Across nodes, operated on in parallel

- Two ways to create them

- *Parallelized Collections*

```
data = [1, 2, 3, 4, 5]  
distData = sc.parallelize(data)
```

- *External Datasets* - Any Hadoop supported storage - HDFS, S3, HBase...

```
distFile = sc.textFile("data.txt")
```

# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

- **Immutable** - Rebuild lost data
- **Fault Tolerant** - Recover data using Lineage

- **Distributed** - Across nodes, operated on in parallel

- Two ways to create them

- *Parallelized Collections*

```
data = [1, 2, 3, 4, 5]  
distData = sc.parallelize(data)
```

- *External Datasets* - Any Hadoop supported storage - HDFS, S3, HBase...

```
distFile = sc.textFile("data.txt")
```

- **Partitions** - How is data distributed, determined by number of cores  
`sc.parallelize(data, 10)`

# RDD

## RDD - Resilient Distributed Dataset

- **Resilient**

- **Immutable** - Rebuild lost data
- **Fault Tolerant** - Recover data using Lineage

- **Distributed** - Across nodes, operated on in parallel

- Two ways to create them

- *Parallelized Collections*

```
data = [1, 2, 3, 4, 5]  
distData = sc.parallelize(data)
```

- *External Datasets* - Any Hadoop supported storage - HDFS, S3, HBase...

```
distFile = sc.textFile("data.txt")
```

- **Partitions** - How is data distributed, determined by number of cores  
`sc.parallelize(data, 10)`

- Generally unstructured data

# RDD

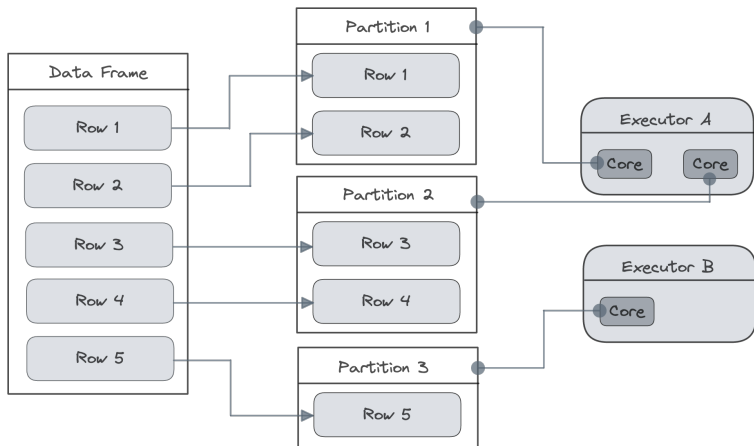


Image 1 - Partitions



# RDD

RDDs support two type of operations:

- **Transformations** - create a new dataset from an existing one i.e. creates a completely new RDD

# RDD

RDDs support two type of operations:

- **Transformations** - create a new dataset from an existing one i.e. creates a completely new RDD
  - `map()` - applies function to each element of a dataset and returns new RDD

# RDD

RDDs support two type of operations:

- **Transformations** - create a new dataset from an existing one i.e. creates a completely new RDD
  - `map()` - applies function to each element of a dataset and returns new RDD
  - All transformations are lazy - think DAG
- **Actions** - return a value to the driver program after running a computation on the dataset

# RDD

RDDs support two type of operations:

- **Transformations** - create a new dataset from an existing one i.e. creates a completely new RDD
  - `map()` - applies function to each element of a dataset and returns new RDD
  - All transformations are lazy - think DAG
- **Actions** - return a value to the driver program after running a computation on the dataset
  - `reduce()` - aggregates all the elements using some function

# RDD

RDDs support two type of operations:

- **Transformations** - create a new dataset from an existing one i.e. creates a completely new RDD
  - `map()` - applies function to each element of a dataset and returns new RDD
  - All transformations are lazy - think DAG
- **Actions** - return a value to the driver program after running a computation on the dataset
  - `reduce()` - aggregates all the elements using some function
  - Transformations are invoked when an action is triggered

# RDD

RDDs support two type of operations:

- **Transformations** - create a new dataset from an existing one i.e. creates a completely new RDD
  - `map()` - applies function to each element of a dataset and returns new RDD
  - All transformations are lazy - think DAG
- **Actions** - return a value to the driver program after running a computation on the dataset
  - `reduce()` - aggregates all the elements using some function
  - Transformations are invoked when an action is triggered

## Note

Directly programming on RDD level is not efficient, introduces latency, lacks control, and is generally discouraged. Use it if you really have to!

# RDD

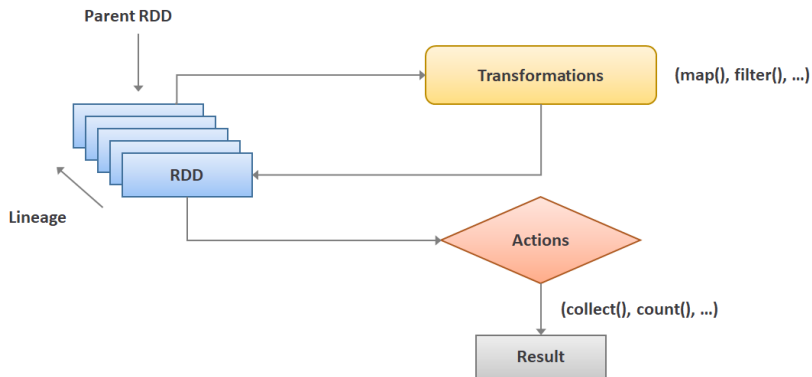


Image 2 - RDD Operations

# Spark Streaming



Image 3 - Spark Streaming



# Spark Streaming



Image 3 - Spark Streaming

- Mini batching - near-real time

# Spark Streaming



Image 3 - Spark Streaming

- Mini batching - near-real time
- Legacy - there are better options now

# Spark Streaming



Image 3 - Spark Streaming

- Mini batching - near-real time
- Legacy - there are better options now
- Abstract concept - **DStream** - Discretized Stream

# Spark Streaming



Image 3 - Spark Streaming

- Mini batching - near-real time
- Legacy - there are better options now
- Abstract concept - **DStream** - Discretized Stream
  - Continuous stream of data

# Spark Streaming



Image 3 - Spark Streaming

- Mini batching - near-real time
- Legacy - there are better options now
- Abstract concept - **DStream** - Discretized Stream
  - Continuous stream of data
  - Represents a continuous sequence of RDDs

# Spark Streaming



Image 3 - Spark Streaming

- Mini batching - near-real time
- Legacy - there are better options now
- Abstract concept - **DStream** - Discretized Stream
  - Continuous stream of data
  - Represents a continuous sequence of RDDs
  - Operations on DStream translates to operations on underlying RDDs




Image 4 - DStream RDDs

# Spark Streaming



Image 3 - Spark Streaming

- Mini batching - near-real time
- Legacy - there are better options now
- Abstract concept - **DStream** - Discretized Stream
  - Continuous stream of data
  - Represents a continuous sequence of RDDs
  - Operations on DStream translates to operations on underlying RDDs




Image 4 - DStream RDDs

# Spark Streaming

- Input DStreams - representing the stream of input data



# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections
  - **Advanced** - Extra utility classes: Kafka, Flume, Kinesis

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections
  - **Advanced** - Extra utility classes: Kafka, Flume, Kinesis
- Drawbacks



# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections
  - **Advanced** - Extra utility classes: Kafka, Flume, Kinesis
- Drawbacks
  - Directly dealing with RDDs

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections
  - **Advanced** - Extra utility classes: Kafka, Flume, Kinesis
- Drawbacks
  - Directly dealing with RDDs
  - Complicated and different API (from batch)

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections
  - **Advanced** - Extra utility classes: Kafka, Flume, Kinesis
- Drawbacks
  - Directly dealing with RDDs
  - Complicated and different API (from batch)
  - Latency

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections
  - **Advanced** - Extra utility classes: Kafka, Flume, Kinesis
- Drawbacks
  - Directly dealing with RDDs
  - Complicated and different API (from batch)
  - Latency
  - Bad Optimization

# Spark Streaming

- Input DStreams - representing the stream of input data
- Associated with **Receiver** - object which receives the data from a source, runs on an allocated core
  - Reliable Receiver - sends ACK to the source
  - Unreliable Receiver - doesn't send ACK
  - We can create custom receivers
- Sources
  - **Basic** - Available in StreamingContext: file systems and socket connections
  - **Advanced** - Extra utility classes: Kafka, Flume, Kinesis
- Drawbacks
  - Directly dealing with RDDs
  - Complicated and different API (from batch)
  - Latency
  - Bad Optimization
  - SQL not supported

# Spark Streaming

- DStream can be combined and joined

# Spark Streaming

- DStream can be combined and joined
- For example, Window operations

# Spark Streaming

- DStream can be combined and joined
- For example, Window operations




Image 5 - DStream Window



# Spark Streaming

- DStream can be combined and joined
- For example, Window operations




Image 5 - DStream Window

However, DStream is not great for Stateful streams and can't handle late data. There are other alternatives!

# DataFrame/DataSet

- Introduced in Spark 2.0

# DataFrame/DataSet

- Introduced in Spark 2.0
- Built on top of RDDs

# DataFrame/DataSet

- Introduced in Spark 2.0
- Built on top of RDDs
- Imposes a structure over data - columns/rows

# DataFrame/DataSet

- Introduced in Spark 2.0
- Built on top of RDDs
- Imposes a structure over data - columns/rows
- Dataset has two APIs - strongly-typed and untyped

# DataFrame/DataSet

- Introduced in Spark 2.0
- Built on top of RDDs
- Imposes a structure over data - columns/rows
- Dataset has two APIs - strongly-typed and untyped
- Row - generic untyped JVM object

# DataFrame/DataSet

- Introduced in Spark 2.0
- Built on top of RDDs
- Imposes a structure over data - columns/rows
- Dataset has two APIs - strongly-typed and untyped
- Row - generic untyped JVM object

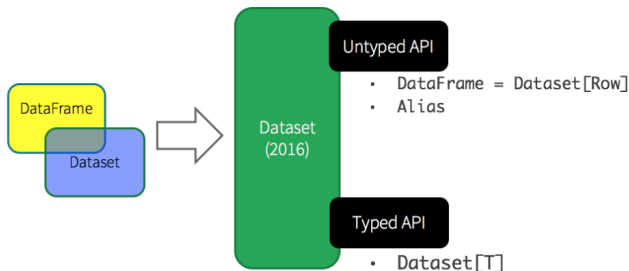


Image 6 - DataFrame/DataSet

# DataFrame/DataSet

- Introduced in Spark 2.0
- Built on top of RDDs
- Imposes a structure over data - columns/rows
- Dataset has two APIs - strongly-typed and untyped
- Row - generic untyped JVM object

Language	Main Abstraction
Scala	Dataset[T] & DataFrame (alias for Dataset[Row])
Java	Dataset[T]
Python	DataFrame
R	DataFrame



# DataFrame/DataSet

## Structure Benefits:

- Syntax check, static-typing, runtime type-safety i.e. catch errors at compile time

# DataFrame/DataSet

## Structure Benefits:

- Syntax check, static-typing, runtime type-safety i.e. catch errors at compile time
- Simpler high-level operations - `agg`, `filter`, `groupBy`...

# DataFrame/DataSet

## Structure Benefits:

- Syntax check, static-typing, runtime type-safety i.e. catch errors at compile time
- Simpler high-level operations - `agg`, `filter`, `groupBy`...
- Uses Spark SQL - has **Catalyst Optimizer** which generates optimized logical and physical plan, together with **Tungstan**

# DataFrame/DataSet

## Structure Benefits:

- Syntax check, static-typing, runtime type-safety i.e. catch errors at compile time
- Simpler high-level operations - `agg`, `filter`, `groupBy`...
- Uses Spark SQL - has **Catalyst Optimizer** which generates optimized logical and physical plan, together with **Tungstan**




Image 7 - Catalyst Optimizer

# DataFrame/DataSet

## Structure Benefits:

- Syntax check, static-typing, runtime type-safety i.e. catch errors at compile time
- Simpler high-level operations - `agg`, `filter`, `groupBy`...
- Uses Spark SQL - has **Catalyst Optimizer** which generates optimized logical and physical plan, together with **Tungstan**




Image 8 - Memory Usage

# Structured Streaming

Builds on top of structure provided by DataFrames/DataSets, which has its benefits:

- Fast and scalable

# Structured Streaming

Builds on top of structure provided by DataFrames/DataSets, which has its benefits:

- Fast and scalable
- Fault-tolerant

# Structured Streaming

Builds on top of structure provided by DataFrames/DataSets, which has its benefits:

- Fast and scalable
- Fault-tolerant
- End-to-end exactly-once stream



# Structured Streaming

Builds on top of structure provided by DataFrames/DataSets, which has its benefits:

- Fast and scalable
- Fault-tolerant
- End-to-end exactly-once stream
- Stream and Batch are the same, we develop once

# Structured Streaming

Builds on top of structure provided by DataFrames/DataSets, which has its benefits:

- Fast and scalable
- Fault-tolerant
- End-to-end exactly-once stream
- Stream and Batch are the same, we develop once
- Provides Event Time handling and Late Data processing - Watermarking

# Structured Streaming

Builds on top of structure provided by DataFrames/DataSets, which has its benefits:

- Fast and scalable
- Fault-tolerant
- End-to-end exactly-once stream
- Stream and Batch are the same, we develop once
- Provides Event Time handling and Late Data processing - Watermarking
- Perform complex SQL queries

# Structured Streaming

How it works:

- Unbounded Input Table

Data stream



Unbounded Table





new data in the  
data stream

=

new rows appended  
to a unbounded table

Image 9 - Unbounded Table

# Structured Streaming

How it works:

- Result Table




Image 10 - Result Table

# Structured Streaming

- Output Modes - defines what gets written to the external storage

# Structured Streaming

- Output Modes - defines what gets written to the external storage
  - **Complete** - Complete result tables are written

# Structured Streaming

- Output Modes - defines what gets written to the external storage
  - **Complete** - Complete result tables are written
  - **Append** - Only rows that get appended since the last trigger



# Structured Streaming

- Output Modes - defines what gets written to the external storage
  - **Complete** - Complete result tables are written
  - **Append** - Only rows that get appended since the last trigger
  - **Update** - Only rows that were updated

# Structured Streaming

- **Output Modes** - defines what gets written to the external storage
  - **Complete** - Complete result tables are written
  - **Append** - Only rows that get appended since the last trigger
  - **Update** - Only rows that were updated

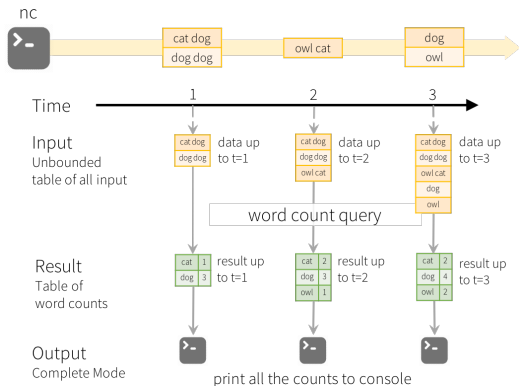


Image 11 - Query Example

# Structured Streaming

## Window operations

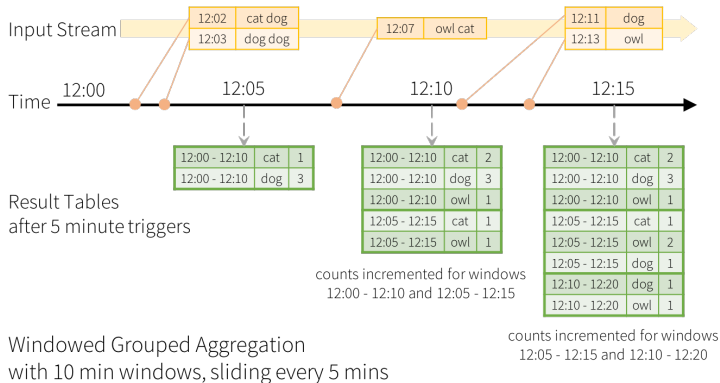


Image 12 - Window Operation

# Watermarking

- Introduced in Spark 2.1

# Watermarking

- Introduced in Spark 2.1
- Keeps track of event time in data

# Watermarking

- Introduced in Spark 2.1
- Keeps track of event time in data
- Engine maintains state and allow late data for specific window  $T$

# Watermarking

- Introduced in Spark 2.1
- Keeps track of event time in data
- Engine maintains state and allow late data for specific window  $T$
- Regulates state - cleaning

# Watermarking

- Introduced in Spark 2.1
- Keeps track of event time in data
- Engine maintains state and allow late data for specific window T
- Regulates state - cleaning
- Example:

```
words = ... # streaming DataFrame

# Group the data by window
windowedCounts = words \
    .withWatermark("timestamp", "10 minutes") \
    .groupBy(
        window(words.timestamp, "10 minutes", "5 minutes"),
        words.word) \
    .count()
```



# Watermark

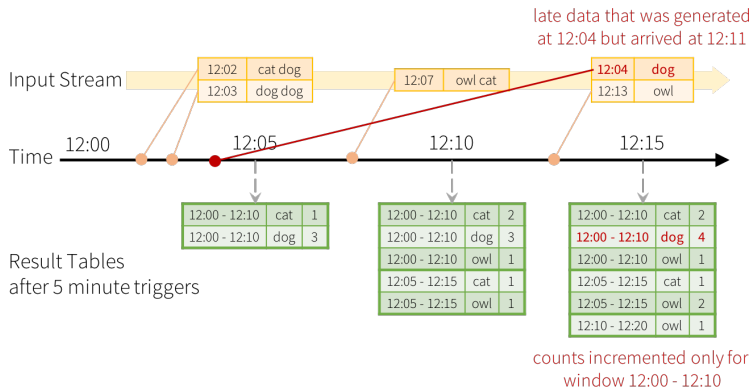


Image 13 - Watermark

# THANKS!

If you have any questions, please don't hesitate to ask

# References

- Spark Official Documentation
  - [RDD Programming Guide](#)
  - [Streaming Programming Guide](#)
  - [Structured Streaming Programming Guide](#)
- Databricks
  - [A Tale of Three Apache Spark APIs: RDDs vs DataFrames and Datasets](#)
  - [Multiple Stateful Operators in Structured Streaming](#)

## Images

- Image 1 - Partitions - [Source](#)
- Image 2 - RDD Operations - [Source](#)
- Image 3 - Spark Streaming - [Source](#)
- Image 4 - DStream RDDs - [Source](#)
- Image 5 - DStream Window - [Source](#)
- Image 6 - DataFrame/DataSet - [Source](#)
- Image 7 - Catalyst Optimizer - [Source](#)
- Image 8 - Memory Usage - [Source](#)
- Image 9 - Unbounded Table - [Source](#)
- Image 10 - Result Table - [Source](#)
- Image 11 - Query Example - [Source](#)
- Image 12 - Window Operation - [Source](#)
- Image 13 - Watermark - [Source](#)